

# Work-in-Progress: Evaluating Stage-Prioritized IDK Cascades in Embedded Real Time Systems

Comyar Dehlavi  
Memorial High School  
Houston, TX, 77024, USA  
dehlavi.comyar@outlook.com

Albert M. K. Cheng  
Department of Computer Science  
University of Houston  
Houston, TX, 77004, USA

Thomas Carroll  
Department of Computer Science  
University of Houston  
Houston, TX, 77004, USA

**Abstract**—IDK (I Don't Know) cascades are classification frameworks that selectively switch between models based on confidence thresholds; when confidence is insufficient, the system abstains ("I don't know") and defers to a deterministic fallback, such as a PID controller. This paper proposes and evaluates a stage-prioritized variant of the IDK cascade, in which constituent neural networks are trained and selected for performance in specific phases of system behavior. The system was implemented on a Raspberry Pi Zero W to regulate the temperature of a Peltier element at a constant 16°C. Two neural networks each with two layers were trained using PID-derived trial data in order to optimize the models with the goals of lowering latency, reducing power usage, and lowering the standard deviation of the temperature readings. The first, faster network (13 neurons) was prioritized during the early cooling phase when the temperature was farthest from the setpoint; the second (49 neurons), more accurate network was favored near steady-state conditions. Results show that this stage-aware cascade outperformed both neural networks individually and the deterministic PID fallback in terms of both temperature stability and power efficiency at high enough confidence thresholds.

## I. INTRODUCTION

In embedded real time systems, striking a balance between lowering computational cost and increasing accuracy of the system is often a major challenge for developers. Factors like power consumption, execution time, and efficacy of the system's output all determine the overall desirability of the system. This often leads to a dichotomy between engineers and developers as to whether they should prioritize well established algorithms or resort to more modern and advanced models such as neural networks for control systems. In recent years, new research has presented IDK (I Don't Know) classifiers which dynamically switch between deep learning models based on a set confidence threshold, in which the system resorts to a reliable deterministic fallback algorithm when the confidence is below the threshold [1]. IDK cascades accomplish the goal of balancing low execution times and relatively high accuracy, making it desirable for embedded systems where both must be accomplished for overall performance of the system [3].

Despite the inherent benefits associated with IDK cascades in embedded systems, they often are underdeveloped in the way they determine which classifier may be best suited depending on the specific stage or segment of the application's

runtime. This issue stems from the iterative nature of most IDK cascades, in which they only switch classifiers and resort to other models with low confidence being the only factor. In embedded systems with a clear runtime course, having the IDK cascade inherently recognize the strengths of each model and giving preference depending on the stage can greatly improve and strengthen the applications overall runtime performance.

When designing an IDK cascade for embedded systems, not only can stage aware processing improve performance [2], but giving individual neural networks awareness of the important factors like system latency and power consumption in their training can work to greatly strengthen the balance of performance-accuracy which embedded developers are actively seeking. Individual networks can then be further refined and specified for their portion of the applications runtime by varying the amount of neurons to control their speed and responsiveness. This allows for networks to be better adjusted for critical portions of the application.

In this work, we evaluated this stage-prioritized and specified implementation of the IDK cascade in embedded systems by applying it to a Raspberry Pi Zero W which was designed to regulate and maintain the temperature of a Peltier element at a baseline of 16°C. Trial data was collected consisting of the average power consumption of the system, standard deviation of the temperature, and average latency. The system consisted of two neural networks along and a PID deterministic fallback, all of which outputted a linear PWM duty cycle to control the extent to which the Peltier element would cool in hopes of regulating its temperature based on sensor readings.

Of the two neural networks, one was faster with 13 total neurons, making it more responsive and preferable for the initial segment of the program where the program must cool more aggressively as the temperature difference is the largest. The second neural network had 49 neurons and was preferred when the read temperature was near the setpoint due to the model's precision. Trial data was collected at varying confidence thresholds, as well as for networks and PID models by themselves.

## II. RELATED WORK

The concept of IDK Cascades as confidence-based classifiers was introduced by Sanjoy Baruah (2023) for switching

based on probabilistic scheduling and having fallback mechanisms which are deadline-aware [1]. While this in theory led to better average latency and performance [2], in the context of embedded systems it was agnostic to the stage of runtime. This paper builds upon that foundation by implementing a version of the IDK cascade that prioritizes certain models based on the portion of the program in which it is currently in, as well as having each neural network be more so targeted for specific parts of the program in addition to also considering hardware limitations and factors like power consumption.

### III. METHODOLOGY

#### A. System Architecture

A Raspberry Pi Zero W was used for implementing the thermal regulation system, which controlled a Peltier thermoelectric cooling element using PWM (Pulse-width modulation). Four TMP36 temperature sensors were attached to the surface of the element, and an average reading was fed into the control algorithm at a fixed sampling interval. The controller's output is a duty cycle between 0 and 100, which determines the intensity of the Peltier cooling via PWM. The system is powered through a regulated 5V supply, and temperature feedback is used to maintain a baseline of 16°C. A schematic of the control circuit is shown in Fig. 1. Additional documentation can be found at this paper's *GitHub* repository.<sup>1</sup>

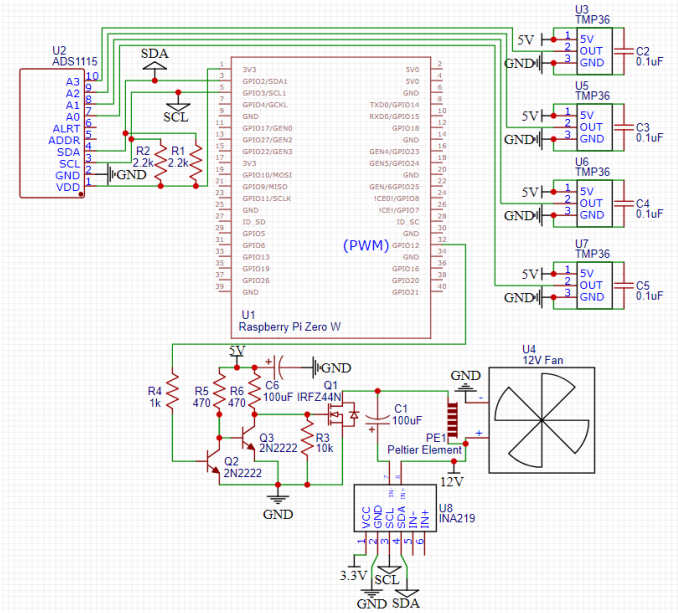


Fig. 1. Circuit schematic of the thermal regulation system using Raspberry Pi Zero W, temperature sensors, and Peltier element.

#### B. IDK Cascade Logic

The controller uses an IDK cascade that switches between three control models: a fast but less precise neural network (NN1), a slower but more accurate neural network (NN2), and a fallback PID controller [3]. At each time step, the controller evaluates the confidence of the predicted duty cycle from a model. If the confidence falls below a fixed threshold, the controller abstains from using the model and activates a fallback, usually the PID controller.

Unlike conventional IDK implementations, which use only confidence to determine switching, this cascade also incorporates a stage-prioritization strategy. NN1 is preferred during the early cooling stage when the system's temperature is far from the setpoint due to its more aggressive control and speed. NN2 is preferred when the temperature is near the setpoint, in which it exhibits greater stability. The PID controller is invoked when the confidence of both networks is too low or ineffective beyond the neural network's capacities.

The IDK cascade also has additional features like hysteresis margins, dwell counters, and a confidence history mechanism. These were included to reduce controller oscillation and ensure stable transitions between models. There was also a deadline of 1.5 seconds so that the models could maintain real-time responsiveness.

#### C. Neural Network Design and Training

The two neural networks were trained using TensorFlow and implemented on the Raspberry Pi Zero W using TensorFlow Lite Micro due to the hardware restraints. Both networks have two hidden layers and a single linear output neuron, but differ in complexity and as a result in their overall strengths.

- **NN1 (Fast Network):** 8 neurons in the first layer, 4 in the second. Total: 13 neurons.
- **NN2 (Slow Network):** 32 neurons in the first layer, 16 in the second. Total: 49 neurons.

Each network receives three input features: current temperature, current power level, and estimated latency of the previous control step. The output is a predicted PWM duty cycle from 0-100. Both models were trained on labeled datasets generated from standalone PID trials, in which the system was controlled using a PID algorithm and data was logged every second for several minutes. The objective during training was to minimize a weighted loss function prioritizing low latency, low power consumption, and minimal standard deviation in temperature.

The fewer number of total neurons in NN1 give it faster processing speed, but also leads to one of its major attributes being more aggressive cooling and outputting a higher duty cycle, which is preferred when the temperature difference between measured and setpoint is largest. Since NN2 has much more neurons, it is thus slower, but more precise and accurate, making it better suited for small fluctuations and when the measured temperature is hovering around the setpoint.

#### D. Stage-Prioritization Strategy

A major feature of this implementation of the IDK-Cascade is the stage-aware prioritization of classifiers. Instead of treating all models the same and selecting solely off of confidence scores, the cascade looks at the difference in recorded and wanted temperatures and gives preference to a model based off that, making it more likely to be selected. The stages are thus determined by the current recorded temperature.

When the recorded temperature is significantly above the chosen temperature, 16°C in our case, the controller favors NN1 due to its responsiveness. As the system approaches the target temperature, or within  $\pm 4^\circ\text{C}$ , and the cooling

<sup>1</sup>All source code, trained models, trial data, and documentation are available at <https://github.com/Void-Overflow/Neural-Network-vs-PID-in-IDK-Cascade-for-Thermal-Regulation-Research>

demand decreases, NN2 becomes the preferred model due to its ability to provide finer control. The PID controller is employed primarily for safety and recovery—especially in the event of unexpected confidence drops, persistent oscillations, or inability for temperature to approach baseline.

#### IV. EXPERIMENTS AND RESULTS

##### A. Experimental Setup

The system was evaluated in a 26°C room environment. Each trial was precisely began with the Peltier surface at room temperature and ran for about 45 minutes. Each trial had a 16°C baseline. Four configurations were tested: PID-only, NN1-only, NN2-only, and the IDK Cascade. Additionally, the IDK cascade was evaluated at three confidence intervals: 30%, 50%, and 70%. The main factors analyzed in each trial were average power consumption, average latency (execution time), temperature stability (standard deviation), and the amount of time each controller was used within the cascade.

##### B. Controller Performance Comparison

Table 1 displays the average power consumption, latency, and accuracy relative to the baseline of each configuration. The PID controller was found to be the fastest model by far in terms of computation speed in comparison to the others at just 3.430 ms as well as a relatively high accuracy of 97.39%. When comparing the standalone neural networks, NN2 was much more accurate than NN1, with 94.781% accuracy in comparison to 83.32%, which is consistent with the training implementation as the majority of the runtime is when the measured temperature is hovering around the baseline, with NN1 being better at the initial phase and NN2 under performing initially. NN1 also is seen with a higher power consumption of any other model at 18.899 W, which is also consistent with the notion of it being a more aggressive coolant and model.

However, when combining the models into IDK cascades, we see that at 50% confidence the cascade performs nearly as well as the fallback PID algorithm, and at 70% the IDK cascade manages to outperform all other models, including the PID, with 98% accuracy and 14 W power consumption. These results suggest that blending NN1’s speed and NN2’s precision through dynamic switching at high enough confidence intervals leads to superior performance.

TABLE I

COMPARISON OF AVERAGE POWER USAGE, LATENCY, AND TEMPERATURE STABILITY ACROSS CONTROL STRATEGIES.

Model	Accuracy (%)	Latency (ms)	Power (W)
PID	97.39	3.430	16.396
NN1	83.32	12.569	18.899
NN2	94.71	12.708	17.324
IDK (30%)	96.44	13.776	15.722
IDK (50%)	97.27	13.713	17.208
IDK (70%)	98.48	13.378	14.225

##### C. Stage Activation Breakdown

Figure 2 shows controller usage distribution within the IDK Cascade at different confidence thresholds. At 30%, we see NN2 being significantly dominant over NN1 and NN2 due to its inherent preference by the cascade since it’s designed stage makes up most of the runtime, but also because it has a very wide ranger of operation with due to the low confidence interval. At 50% confidence we see a slightly lower portion of the program being dominated by NN2, but we also see the PID and NN1 controller be active about the same percentage of the time. At 70% confidence, which is when we notice the cascade start to outperform all other metrics, we notice that the PID is being executed the vast majority of the runtime, while NN2 is only slightly higher than NN1. This is consistent with the confidence threshold being high as the networks are more limited and more often output "I don’t know."

The stage-prioritization implemented alongside the traditional IDK cascade framework leads to a more predictive distribution of models and allows the developer to exhibit greater control on how the cascade specifically interacts and influences the output and model selection.




Fig. 2. Proportion of controller usage at varying IDK confidence thresholds.

##### D. Temperature Regulation Over Time

Figure 3 depicts the temperature of each control model vs time over a major portion of each 45 minute trial. A significant finding can be that NN1 cools to a much greater extent compared to the other models and much faster; however, it overshoots past the 16°C baseline and continued to overshoot, being approximately 12.5°C towards the end of the trial. In the NN2 trial, significant osculations can be observed in comparison to the IDK cascade models, along with a recognizable amount of oscillation towards the middle of the trial for the IDK cascade at a 30% confidence threshold. The PID, 50% IDK, and 70% IDK experienced minimal osculations and overall tended to stick much better to the 16°C baseline.

These traces demonstrate that combining these models and focusing on their strengths during certain portions in moderation leads ultimately to the most stable possible execution in terms of thermo-efficiency.

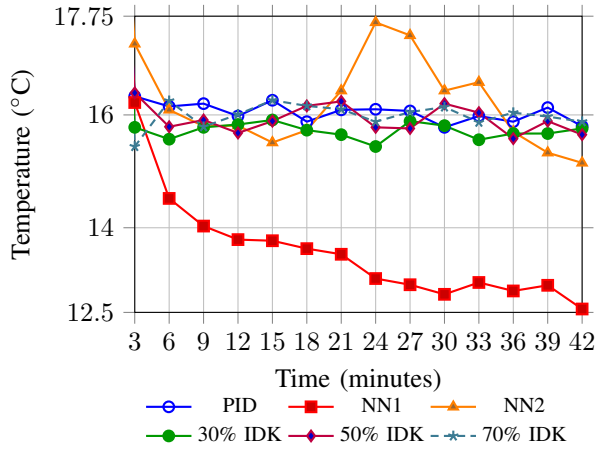


Fig. 3. Temperature readings for different controllers with baseline at 16°C.

## V. DISCUSSION

The results of this publication display strong results that the stage-aware implementation of the IDK cascade can substantially improve the performance of embedded systems. Traditional IDK models make switching decisions based purely off model confidence, which isn't as efficient in the context of embedded systems with additional hardware and runtime limitations which are needed to be considered [3]. However, despite the additional challenges associated with embedded systems, they often follow predictable behavioral stages (e.g., initialization, stabilization), and actively following and accommodating these runtime behaviors in model decision can lead to considerable performance gains. By distinguishing certain neural networks for certain predictable phases of the runtime, and allowing a fallback algorithm in times of uncertainty, this allows the IDK cascade to be much more precise and not only more accurate in its output, but also more efficient in terms of using limited computational resources.

The performance gains seen in this experiment are multifaceted. The stage-aware IDK cascade at 70% confidence not only managed to increase the accuracy and therefore stability of the trial to be higher than that of the fallback, but it also reduced power consumption when compared to any of the other models. This highly suggests that utilizing preference for certain networks based on the segment of runtime which it is in, as well as specifically training the networks to be aware of hardware limitations, directly addresses the issue of balancing responsiveness and stability with energy efficiency in embedded environments. Furthermore, the application of hysteresis and threshold tuning helped prevent oscillations between models, reduced jitter, and help smooth and make transitions between models more smooth.

Compared to the IDK cascade framework presented by Baruah et al. [1], the implementation presented in this paper introduced a more deliberate selection mechanism that implemented greater awareness of runtime and hardware limitations in systems where computational power is finite

rather than just focusing on confidence scores alone. While the analysis presented in that work focused on optimizing switching primarily on minimizing execution deadlines [2], this paper expands on that work by considering factors that are additionally significant in embedded systems, especially when tied to physical processes like thermo-regulation.

## VI. CONCLUSION

The findings of this paper present a further enhanced version of the IDK cascade specifically tailored to the environment of embedded real time systems, in which it combined runtime stage-awareness with the traditional switching associated with IDK classifiers. In the case study of thermal regulation with a Raspberry Pi Zero W and Peltier element, the stage-aware cascade ultimately proved more effective compared to both traditional PID control and standalone neural networks. It had tighter temperature control, lower power usage, and ultimately more control over the runtime.

Aside from just validating the efficacy of stage-prioritized IDK switching, this paper provides a framework for embedded systems where it is the upmost important to maintain a very low computational overhead yet exceptionally precise and stable decision making, in which the control method has to not only factor in times of uncertainty, but learn where they belong in the control flow of the program. This method has practical relevance for many developers who are building embedded systems that must navigate performance constraints while aiming for precision.

The work presented in this article specifically regarding the stage-aware cascade in terms of thermo-regulation can be applied to broader control applications such as motor control, HVAC systems, or robotic actuation. Another area of further insight could be to additionally restructure the cascade over time and adjust preference to certain models in embedded systems where the runtime course is variable and not predictable.

Ultimately, this work contributes to the growing field of hybrid intelligent-deterministic systems. The need for the dynamic, precise, yet efficient stage-aware control model presented in this study is especially relevant in this day-and-age where applications and innovations are using increasingly complex and demanding learning-based models. The stage-prioritized IDK cascade demonstrates that its implementation into systems can ultimately yield meaningful gains in precision and efficiency-insights that could benefit a wide variety of embedded applications in the future.

## REFERENCES

- [1] S. Baruah, A. Burns, R. I. Davis, and Y. Wu, "Optimally ordering IDK classifiers subject to deadlines," *Real-Time Systems*, vol. 59, no. 1, pp. 1–34, 2023. doi: <https://doi.org/10.1007/s11241-022-09383-w>.
- [2] S. Baruah, T. Abdelzaher, K. Agrawal, A. Burns, R. I. Davis, Z. Guo, and Y. Hu, "Scheduling IDK classifiers with arbitrary dependences to minimize the expected time to successful classification," *Real-Time Systems*, vol. 59, pp. 348–407, 2023. doi: <https://doi.org/10.1007/s11241-023-09395-0>.
- [3] X. Wang, Y. Luo, D. Crankshaw, A. Tumanov, F. Yu, and J. E. Gonzalez, "IDK cascades: Fast deep learning by learning not to overthink," *arXiv.org*, 2018. <https://arxiv.org/abs/1706.00885>.